

# An Introduction to Large Language Models (LLMs)

Brent McPherson, PhD  
2024-05-17



**McGill**  
UNIVERSITY

# Topics

- What are Large Language Models?
- The Impact of LLMs
- Using an LLM Programmatically
  - Some different tools
  - Retrieval Augmented Generation
- What models should we use?

Hello?

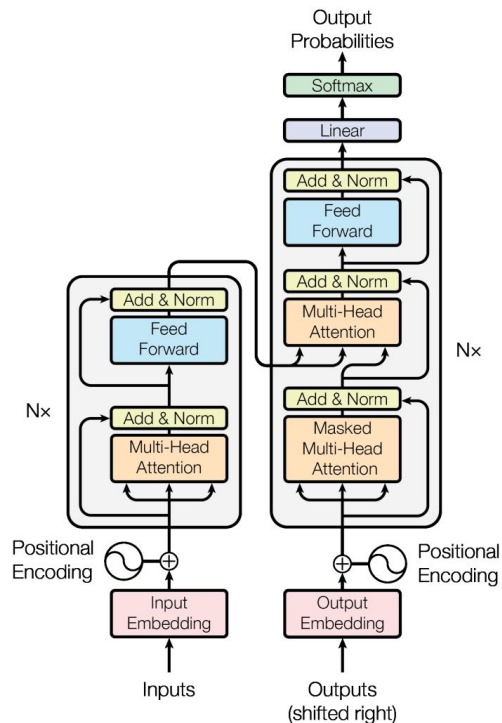


■ ■ ■

# What are Large Language Models?

# What are Large Language Models (LLMs)?

- Massive (billions of parameters) deep learning models.
  - Mostly transformers (encoder / decoder)
- They are trained on a massive corpus (trillions of words) to probabilistically generate responses.
- Conventionally, they provide a plain language interface to interact with them.
  - No “coding” required



# The Parts of an LLM

## ❖ Tokenization

- splitting the input / output texts into smaller units that can be processed by the LLMs.

## ❖ Embedding

- a high dimensional encoding of tokens that represents their semantic meaning.

## ❖ Attention

- LLMs selectively weight the importance of words within the context of their embeddings.

## ❖ Pre-training

- unsupervised training stage of an LLM on large amounts of text to establish word embeddings.


## ❖ Transfer Learning

- utilize the pre-trained weights to provide context to more specialized or fine-tuned data.


# What are Large Language Models (LLMs)?

- A transformer model **embeds** (encodes) the **data** in a high dimensional space.
- To generate responses, the prompt is parsed into that **embedding** space to identify **relevant information** that is **based on weights from the training data**.
- The kinds of responses are **fine-tuned** with more targeted kinds of training inputs.
- The response is generated by providing the next most likely segment (**token**) of text.

Tokenization - Embedding - Attention - Pre-Training - Transfer Learning




# Building an LLM



Embedded  
with **Attention**

An orange arrow pointing right, containing the text "Embedded with Attention".



Into the  
Pre-trained  
Language  
Model

Text describing the step of embedding data into the pre-trained language model.

Fine-tune  
Prompts

A blue arrow pointing right, containing the text "Fine-tune Prompts".

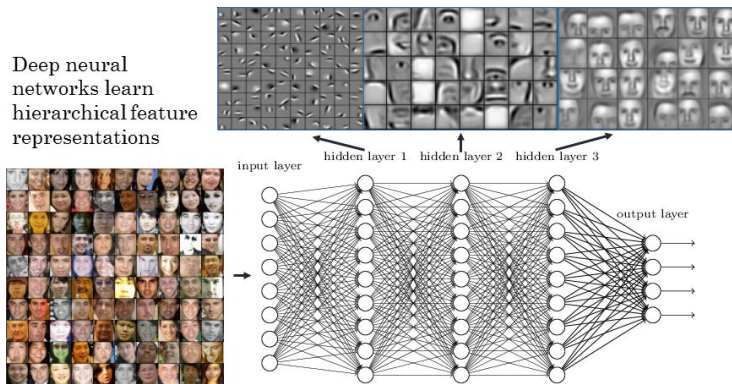


GPT-4o

Tokenization - Embedding - Attention - Pre-Training - Transfer Learning

# What are Large Language Models (LLMs)?

- During model **training**, all the text is **encoded** into the model space as **tokens**.
  - A **token** can be letters and punctuation, but it can also be whole words or parts of speech.
- The layers of the LLM weight the **tokens** together into more complex topics.
  - Similar to the layers in vision networks, segmenting image facets of increasing complexity.
- The further apart phrases or topics are in the **trained semantic space**, the less related they are.





# What are Large Language Models (LLMs)?

- The **prompt is evaluated** in the **embedding** space to determine what parts are most **important to attend** to.
- The model can be **augmented** or **specialized** for different tasks or domains.
  - This is expensive in every sense of the word.
- The **pre-trained** set of **embeddings** can be **transferred** to perform better for different kinds of **specialized prompts**.
  - Again, this is similar to the vision parsing networks.

Tokenization - Embedding - Attention - Pre-Training - Transfer Learning

# Questions?



# The Impact of Large Language Models

# How are Large Language Models Used?

- Text Generation (Chatbots)
  - Customer Service
  - [Talk with your documents](#)
- Coding assistants
  - Github Copilot
  - [Coding teams](#)
- News Reports - Summaries
  - [AI generated news stories](#)
  - Academic Work
- Search Engines
  - [Google called a code red](#)
  - Databases - Search Replacement?



# And everyone is hyped!



Code 55% Faster!

**Nvidia CEO predicts the death of coding — Jensen Huang says AI will do the work, so kids don't need to learn**

News

By Benedict Collins published February 26, 2024

Jensen Huang believes coding languages are a thing of the past

## Revolutionizing Data Annotation: The Pivotal Role of Large Language Models

By **Adnan Hassan** - March 3, 2024

## On the Utility of Large Language Model Embeddings for Revolutionizing Semantic Data Harmonization in Alzheimer's and Parkinson's Disease

Yasamin Salimi, Tim Adams, Mehmet Can Ay, Helena Balabin, Marc Jacobs, and 1 more



# And everyone is hyped!... right?



“Downward Pressure on  
Code Quality”



## Use of large language models might affect our cognitive skills

[Richard Heersmink](#) ✉

[Nature Human Behaviour](#) (2024) | [Cite this article](#)



NEWS | 10 April 2024


## Is ChatGPT corrupting peer review? Telltale words hint at AI use

A study of review reports identifies dozens of adjectives that could indicate text written with the help of chatbots.

By [Dalmeet Singh Chawla](#)

# This isn't the first time...

- The industrial revolution greatly reduced needs of common types of skilled labor.
- A period of upheaval and consolidation, but eventually a new normal.
  - Improved quality of life (arguably).
- However, not everybody was on board.
  - Luddites



# How do we do this right?


- There is a lot of potential for LLMs to greatly improve our lives.
  - Automate away many more tedious, detail oriented tasks.
  - Provide easier access to information.
- There is a lot of problems that may arise from their unchecked use.
  - They are still very unreliable - [hallucinations](#).
  - Corporate control of the largest models is driven by profit, not benefit.
- It is critical that we proactively engage with and understand how to best use this technology so it doesn't cause more harm than good.



# Using an LLM Programmatically

# Prompt Engineering

- A recent buzzword for the “job of the future” thanks to LLMs.
- It is about understanding how to interact with LLMs in the most effective way.
  - Without having to retrain them.
- There are many tools to interacting with LLMs to broaden their general utility.





## How To Google It

**site:**

Only searches the pages of that site.

**“ ”**

Searches for the exact phrase, not each of the words separately.

**-**

Excludes this term from the search.

**site:nytimes.com ~college “test scores” -SATs 2008..2010**

**~**

Will also search related words, such as ‘higher education’ and ‘university’.

**..**

Shows all results from within the designated timerange.

# Anatomy of a Prompt

- There are 3 speakers who can “talk” in a prompt.
  - The user
  - The model
  - The system
- Prompts are typically passed as JSON that hold the history of the interaction.

```
import os
from openai import OpenAI

# pull api key from loaded environment variable
client = OpenAI(api_key=os.getenv("OPENAI_API_KEY"))

# build a helper function w/ new API to return messages
def get_completion(prompt, model="gpt-3.5-turbo"):
    messages = [{"role": "user", "content": prompt}]
    response = client.chat.completions.create(messages=messages, model=model)
    return response.choices[0].message.content

# test the call
get_completion("Why is the sky blue?")
```

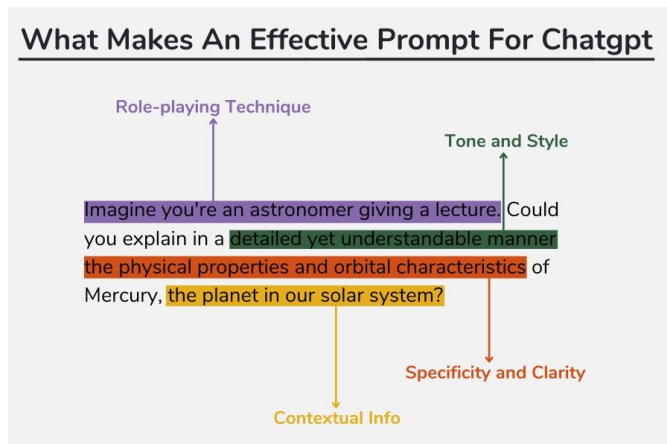
```
def get_completion_from_messages(messages, model="gpt-3.5-turbo", temp=0):
    response = client.chat.completions.create(model=model,
                                              messages=messages,
                                              temperature=temp)
    return response.choices[0].message.content

messages = [
    {'role': 'system', 'content': 'You are an assistant that speaks like Shakespeare.'},
    {'role': 'user', 'content': 'tell me a joke'},
    {'role': 'assistant', 'content': 'Why did the chicken cross the road'},
    {'role': 'user', 'content': 'I don\'t know'} ]

response = get_completion_from_messages(messages, temperature=1)
```

# Anatomy of a Prompt

- The **user** is what you (the user) are asking the model to produce.
- The **model** is the LLMs responses to any input.
- The **system** is a silent setting that can modify how the LLM responds.



# Prompt Engineering - General Guidelines

1. Be clear and specific in what you ask.

- This does not mean be brief.

2. Use delimiters to separate out specific parts or the prompt.

- i.e. summarize the text in triple quotes.
- Helps to sanitize the inputs of the LLMs.

3. Ask for structured output in the response

- JSON formatting, HTML tables, etc.

4. Ask the LLM to check input conditions as part of the prompt.

5. Provide successful input / output pairs for it to mimic.

6. Give the model time to “think”.

- Give the LLM instructions for incremental steps toward the desired outcome
- Have the LLM show its work while hiding its reasoning from the user.

7. Tell the model to make its own solution and compare it to the input

- Again, step by step helps.

```
text = f"""
You should express what you want a model to do by \
providing instructions that are as clear and \
specific as you can possibly make them. \
This will guide the model towards the desired output, \
and reduce the chances of receiving irrelevant \
or incorrect responses. Don't confuse writing a \
clear prompt with writing a short prompt. \
In many cases, longer prompts provide more clarity \
and context for the model, which can lead to \
more detailed and relevant outputs.
"""

prompt = f"""
Summarize the text delimited by triple backticks \
into a single sentence.
```{text}```
"""

response = get_completion(prompt)
print(response)
```

```
prompt = f"""
Generate a list of three made-up book titles along \
with their authors and genres.
Provide them in JSON format with the following keys:
book_id, title, author, genre.
"""

response = get_completion(prompt)
print(response)
```



```

text_1 = f"""
Making a cup of tea is easy! First, you need to get some \
water boiling. While that's happening, \
grab a cup and put a tea bag in it. Once the water is \
hot enough, just pour it over the tea bag. \
Let it sit for a bit so the tea can steep. After a \
few minutes, take out the tea bag. If you \
like, you can add some sugar or milk to taste. \
And that's it! You've got yourself a delicious \
cup of tea to enjoy.
"""

prompt = f"""
You will be provided with text delimited by triple quotes.
If it contains a sequence of instructions, \
re-write those instructions in the following format:

Step 1 - ...
Step 2 - ...
...
Step N - ...

If the text does not contain a sequence of instructions, \
then simply write \"No steps provided.\"

\\\"\\\"\\\"{text_1}\\\"\\\"\\\"
"""

response = get_completion(prompt)
print("Completion for Text 1:")
print(response)

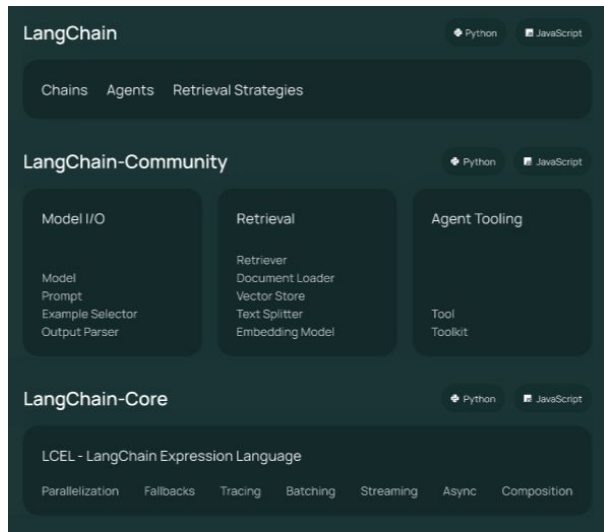
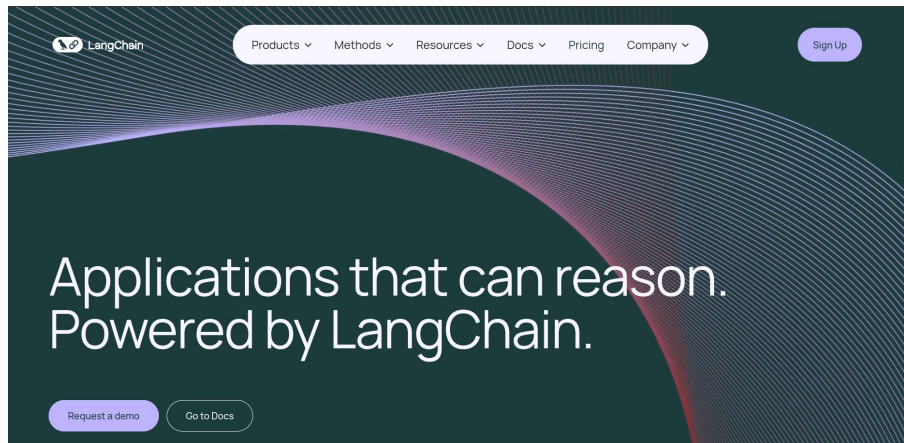
```

# Specific Tools

# LangChain

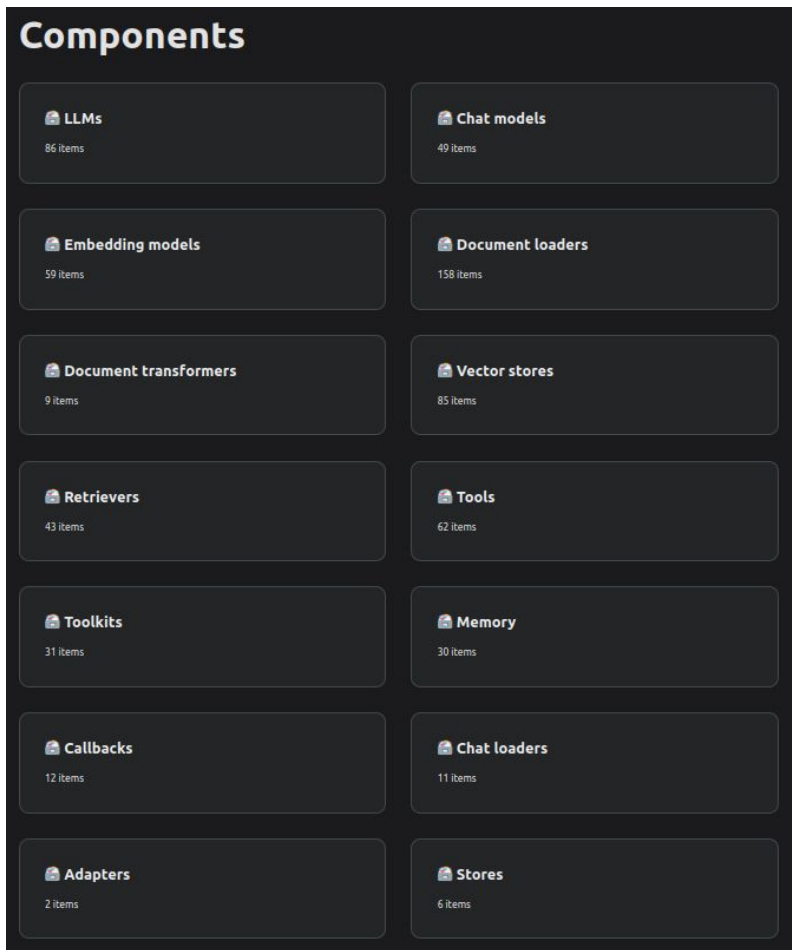


- Python / JS tools for interacting with LLMs programmatically.
- Facilitates prompt engineering and multiple calls to answer a question.
- Many utilities for parsing data for vector embedding in **RAGs**.
- Allows for (relatively) painless swapping between different LLMs.



# LangChain Components

- Classes / Modules for working with standard components.
- Prompt Templates can be created with many standard strategies and logic.
  - LCEL (LangChain Expression Language)
- Because LangChain prompts are modular, they can be easily passed to multiple models.



```
from langchain_core.output_parsers import StrOutputParser
from langchain_core.prompts import ChatPromptTemplate
from langchain_openai import ChatOpenAI



prompt = ChatPromptTemplate.from_template("tell me a short joke about {topic}")
model = ChatOpenAI(model="gpt-4")
output_parser = StrOutputParser()

chain = prompt | model | output_parser

chain.invoke({"topic": "ice cream"})
```



- Many overlapping features with LangChain.
- A bit more focus on building Agents.
- These tools are all still rapidly being developed.
  - Some features / models may be further developed in one ecosystem over another.





## Programming—not prompting—Language Models

- [Newer framework](#) / philosophy.
- Priority on building Agents.
- Optimize your prompts on an LLM without manually adjusting inputs.
  - Rewrite prompts to optimize a cost function to get closer to your desired output.



### Systematic Optimization

Choose from a range of optimizers to enhance your program. Whether it's generating refined instructions, or fine-tuning weights, DSPy's optimizers are engineered to maximize efficiency and effectiveness.



### Modular Approach

With DSPy, you can build your system using predefined modules, replacing intricate prompting techniques with straightforward, effective solutions.



### Cross-LM Compatibility

Whether you're working with powerhouse models like GPT-3.5 or GPT-4, or local models such as T5-base or Llama2-13b, DSPy seamlessly integrates and enhances their performance in your system.

# Scrapegraph.ai

- Combines LLM prompting with web scraping for structured data extraction.
- Web crawlers + LLMs
- Use the flexibility of LLMs parsing to ease the ingestion of data for processing.





# Ollama

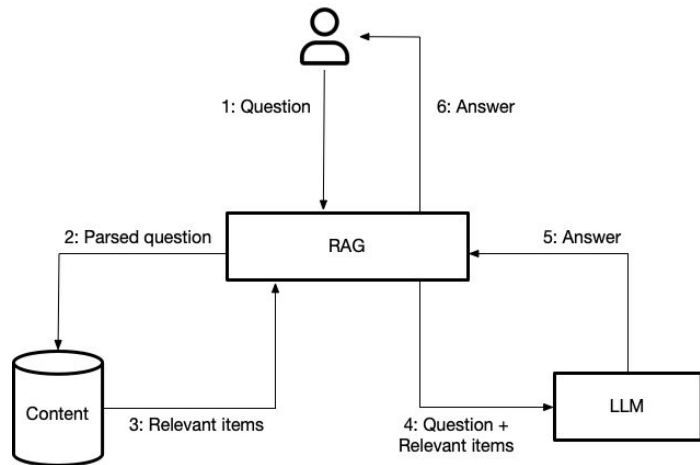
- Locally host and run any number of LLMs.
- Integrates with the previously described tools.
- Designed to be similar to running Docker containers.



What are they all using?

# Retrieval Augmented Generation (RAG)

- It is a standard way of augmenting a model to specialize in a task without having to retrain the model.
- This involves a multistep process.
  - LangChain, LlamaIndex, and many other tools can facilitate this.
- The database (vector store) of information is very important.



```

from langchain_community.vectorstores import DocArrayInMemorySearch
from langchain_core.output_parsers import StrOutputParser
from langchain_core.prompts import ChatPromptTemplate
from langchain_core.runnables import RunnableParallel, RunnablePassthrough
from langchain_openai.chat_models import ChatOpenAI
from langchain_openai.embeddings import OpenAIEmbeddings

vectorstore = DocArrayInMemorySearch.from_texts(
    ["harrison worked at kensho", "bears like to eat honey"],
    embedding=OpenAIEmbeddings(),
)
retriever = vectorstore.as_retriever()

template = """Answer the question based only on the following context:
{context}

Question: {question}
"""

prompt = ChatPromptTemplate.from_template(template)
model = ChatOpenAI()
output_parser = StrOutputParser()

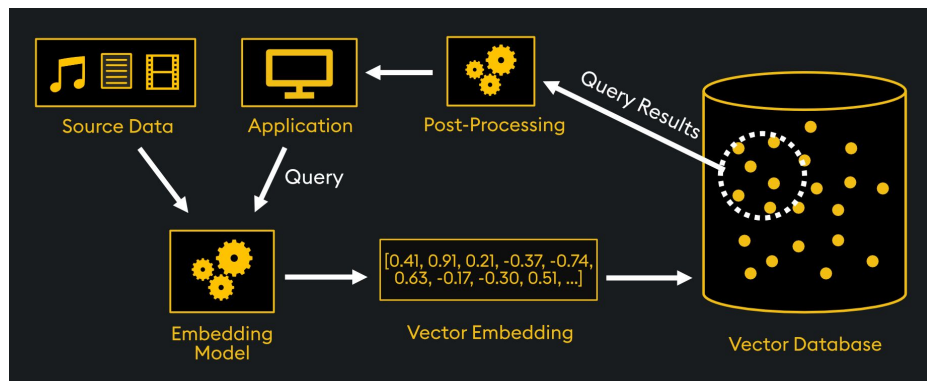
setup_and_retrieval = RunnableParallel(
    {"context": retriever, "question": RunnablePassthrough()}
)
chain = setup_and_retrieval | prompt | model | output_parser

chain.invoke("where did harrison work?")

```

# Vector Databases

- It is a database of information that is embedded in the model's space.
- The prompt is passed to a vector database to identify other relevant information based on the semantic distance between the prompt and the encoded information.
- The identified information is used to augment the prompt and improve the performance of the model.

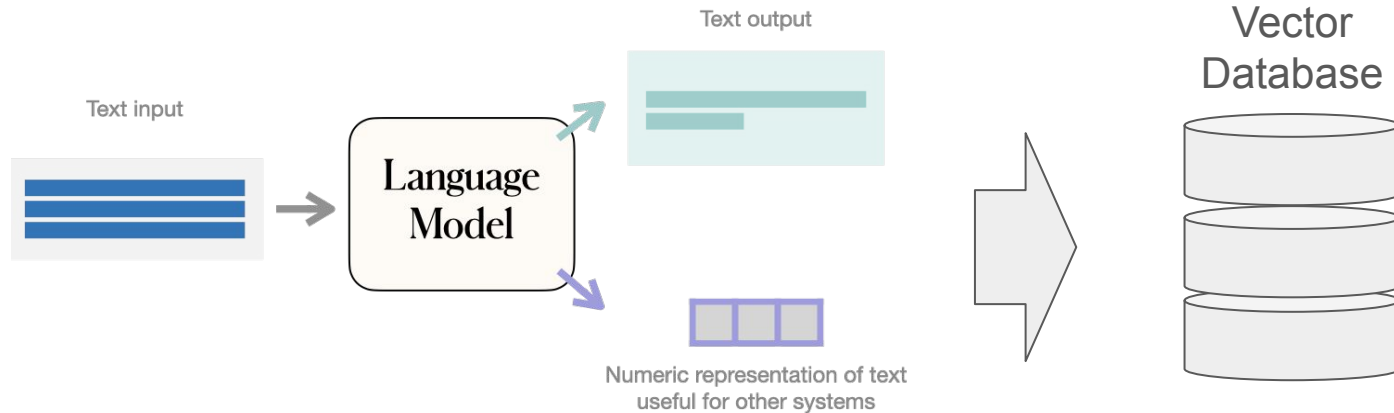


# All kinds of data can be embedded into the vector store

- Pubmed / Arxiv
- Git / Github
- Slack / Gmail
- Project Gutenberg
- Obsidian / Notion
- YouTube transcripts
- Google Services
- AWS
- Microsoft Office Suite
- Most common databases

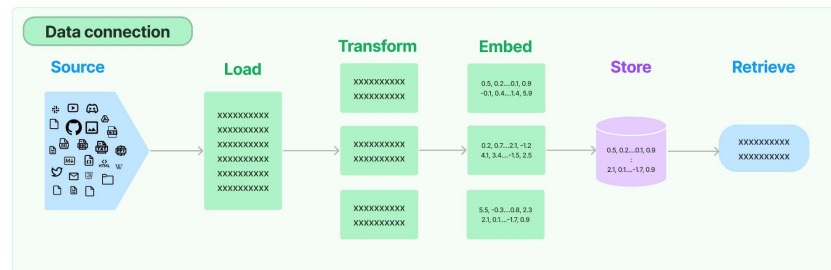
# So what do we need?

- A trained large language model (LLM)
- A way to encode text to the LLMs embedding dimension (Embedding)
- A way to store the text (Vector Database)
- A way to manage the data going between components (LangChain, etc.)



# Preparing Data for a RAG

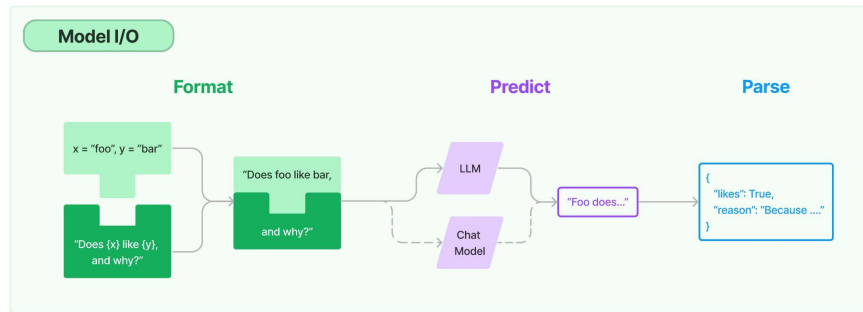
- Extract the text from its source
  - There are many parsers to ease the extraction.
    - Plain text, PDF, YouTube video transcript, GitHub, Notion, Slack, Databases, ...
- Split the text into chunks
  - The chunks have a size and an overlap.
  - The effectiveness with which the data will be encoded can be tuned with these settings.
- Encode the chunks in the model space.
  - Build a vector database of contexts you want available when using the model.





# Pass a Prompt for Context

- Embed the prompt in the vector store.
- Determine the nearby entries based on their distance from each other in the high dimensional space.
  - This can be based on absolute distance or probabilistic variations.
    - Maximum Marginal Variance
- Return that information as text to augment the LLM prompt.

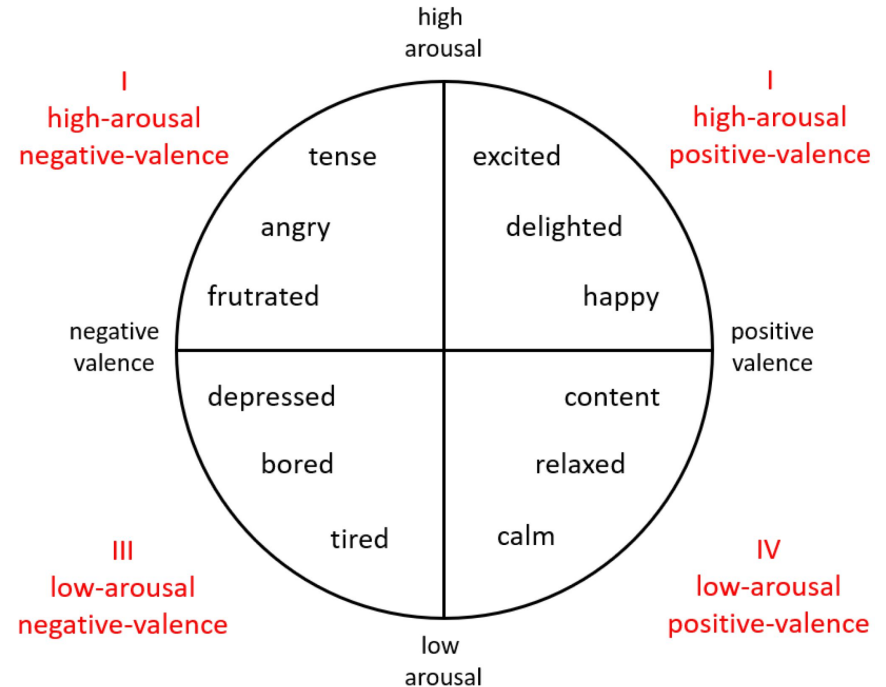


Remember, at any point along the way you could be passing (chaining) your input through another LLM call!

# Specialized LLM Improvements for Intermediary Prompts

- Content Filters / Censors
- Emotional Valence
- Metadata Taggers

These can provide their own metadata about part of a prompt.



What models should we use?

# Open Models. Obviously.


- OpenAI is the most popular
  - Its name is also a lie.
- Google ~~Bard~~ / PaLM / Gemini
  - Costs money.
- Meta / Facebook
  - ~~Llama2~~ Llama3
  - Open sourced.
  - Many specialized derivatives.
- HuggingFace
  - A full repository of open models.



# Conclusion

- LLMs are a new technology that can fundamentally change how we approach our work.
- It is important to proactively engage and understand their functionality and limits.
- There are many emerging (and open) tools to programmatically work with them.
- Start getting creative solving your problems with LLMs.

Hello?



Hello!





## OWNER OF GAMING SITES FIRES WRITERS, HIRES FOR "AI EDITOR" TO CHURN OUT HUNDREDS OF ARTICLES PER WEEK

... culled the jobs of at least 50 humans — annihilating, by [some estimates](#), around 40 percent of its workforce.

... would use AIs like ChatGPT to output up to an astounding — if not outright impossible — 200 to 250 articles of questionable quality *per week*.

Assuming a five-day work week, that's a lot to ask of one person to do with up to 50 articles per day, or less than *ten minutes* per article...

The pay isn't remarkable, either. At the listed salary range of \$40,000 to \$55,000 per year, that works out to roughly \$4.23 per article, **at best**.